

mlr3 Resampling: An Introduction

Natalie Foss

2022-09-19

Introduction

The purpose of this post is to introduce you to `mlr3` resampling and give an example of how to train a model with resampling (showing the best practice methods), using the `palmer penguins` dataset. `palmer penguins` is a built in task described as “classification data to predict the species of penguins”. Read more about the dataset [here](#).

This blog post picks up from the previous post [If you do not wish to read that post](#), this code chunk will get you up to speed (though we recommend you skim it).

```
library("mlr3")
task = tsk("penguins")
learner = lrn("classif.rpart")
measure = msr("classif.ce")
```

Review

In the last post we saw how depending on the random seed we can get performance estimates that vary widely:

```
set.seed(3)
splits = partition(task, ratio = 0.8, cat_col="species")
learner$train(task, splits$train)
prediction = learner$predict(task, splits$test)
prediction$score(measure)
```

```
## classif.ce
## 0.1014493
```

```
set.seed(22)
splits = partition(task, ratio = 0.8, cat_col="species")
learner$train(task, splits$train)
prediction = learner$predict(task, splits$test)
prediction$score(measure)
```

```
## classif.ce
## 0.01449275
```

In order to report the accurate predictive performance of your model, you need resampling.

Resampling methods

Similar to learners and performance measures, there are many popular resampling method implemented. You read about the options here, or just view them:

```
as.data.table(mlr_resamplings)
```

##	key	label	params	iters
## 1:	bootstrap	Bootstrap	ratio, repeats	30
## 2:	custom	Custom Splits		NA
## 3:	custom_cv	Custom Split Cross-Validation		NA
## 4:	cv	Cross-Validation	folds	10
## 5:	holdout	Holdout	ratio	1
## 6:	insample	Insample Resampling		1
## 7:	loo	Leave-One-Out		NA
## 8:	repeated_cv	Repeated Cross-Validation	folds, repeats	100
## 9:	subsampling	Subsampling	ratio, repeats	30

Each resampling method has a set of parameters (0, 1, 2). To view the parameters and other meta-information:

```
rsmp("bootstrap")$param_set
```

```
## <ParamSet>
##      id      class lower upper nlevels      default value
## 1:  ratio ParamDbl      0      1      Inf <NoDefault[3]>      1
## 2:  repeats ParamInt      1      Inf      Inf <NoDefault[3]>     30
```

Cross Validation is the resampling method we will use in this example, read about it here (called k-fold cross validation).

To initialize the resampling object we use the sugar function `rsmp()`. We specify any parameters that we want different than the defaults here.

```
cv = rsmp("cv", folds=10)
```

Now lets perform the resampling. We will set `store_models` to `TRUE` so we can look at individual models later (this defaults to `FALSE` to limit memory consumption).

```
rr = resample(task, learner, cv, store_models = TRUE)
```

```
## INFO [14:32:58.525] [mlr3] Applying learner 'classif.rpart' on task 'penguins' (iter 1/10)
## INFO [14:32:58.576] [mlr3] Applying learner 'classif.rpart' on task 'penguins' (iter 2/10)
## INFO [14:32:58.599] [mlr3] Applying learner 'classif.rpart' on task 'penguins' (iter 3/10)
## INFO [14:32:58.617] [mlr3] Applying learner 'classif.rpart' on task 'penguins' (iter 4/10)
## INFO [14:32:58.633] [mlr3] Applying learner 'classif.rpart' on task 'penguins' (iter 5/10)
## INFO [14:32:58.649] [mlr3] Applying learner 'classif.rpart' on task 'penguins' (iter 6/10)
## INFO [14:32:58.692] [mlr3] Applying learner 'classif.rpart' on task 'penguins' (iter 7/10)
## INFO [14:32:58.711] [mlr3] Applying learner 'classif.rpart' on task 'penguins' (iter 8/10)
## INFO [14:32:58.731] [mlr3] Applying learner 'classif.rpart' on task 'penguins' (iter 9/10)
## INFO [14:32:58.749] [mlr3] Applying learner 'classif.rpart' on task 'penguins' (iter 10/10)
```

```
rr
```

```
## <ResampleResult> of 10 iterations  
## * Task: penguins  
## * Learner: classif.rpart  
## * Warnings: 0 in 0 iterations  
## * Errors: 0 in 0 iterations
```

Now for the moment of truth, we can look at the performance **aggregated** over all resamplings.

```
rr$aggregate(measure)
```

```
## classif.ce  
## 0.05260504
```

Compare this classification error to the ones from the review section (0.1014493 and 0.01449275). It turns out that the true performance is between these values.

Learn about how to use pipelines to make resampling simpler [here](#).

Or learn about benchmarking and how it can help you compare the performance of different learners and different tasks [here](#).